Local number variance as a redistribution handle in the one-dimensional Bose–Hubbard chain

Kunal Bhatia*
Independent Researcher, Heidelberg, Germany†
(Dated: April 29, 2026)

We test whether the local occupation variance $F_i = \langle \hat{n}_i^2 \rangle - \langle \hat{n}_i \rangle^2$ identifies intervention subsets with enhanced finite-horizon response in a one-dimensional open Bose–Hubbard chain under Lindblad dephasing. Additional dephasing at the top-k highest-F_i sites is compared against matched-budget random targeting across 100 independent trials per condition, using exact Lindblad evolution in the fixed-particle-number sector ($L \in \{6, 7, 8\}$, $\tau \in \{1, 2, 3\}$). The primary sweep reveals three regimes: at $J/U = 0.12$ the response is near-zero to weakly harmful (95% CI below zero at the paper's burn-in time, but small in magnitude and partly burn-in-sensitive at $L = 6$); near $J/U \approx 0.20$ there is a crossover that is sensitive to system size and horizon; at $J/U \geq 0.30$ targeted intervention is reliably *beneficial* (CI above zero at every tested (L, τ) combination), with best effects at $J/U = 0.40$, $\tau = 3$: $+0.047$ ($L = 6$), $+0.068$ ($L = 7$), $+0.072$ ($L = 8$). To disentangle variance information from geometry, we perform four additional experiments. A disorder selector sweep ($L = 6$, 50 realizations, 9 selectors) at strong disorder ($\mu_{\max} \geq 1.0$) shows that F_i is the top selector, ranking above geometric centrality (geo), mean occupation (maxn/minn), disorder amplitude (dis_amp), and a generator-action selector (gen); the F_i–geo gap grows monotonically with disorder strength and J/U, confirming that F_i carries information beyond raw spatial position. Shell-matched permutation controls ($L = 6$, $\mu_{\max} \leq 0.3$) show that F_i beats all within-shell permutations by $+0.002$–$+0.006$ at $\tau = 3$, $J/U \geq 0.30$, establishing that the advantage is not reducible to shell geometry. A deterministic inhomogeneous tilt chain ($L = 6$) where $F_i \neq$ geo by construction gives $F_i >$ geo with a gap of $+0.032$–$+0.054$ at $J/U = 0.40$, $\tau = 3$, confirming the result under designed symmetry breaking. A scan of the extra dephasing rate ($\gamma_{\text{extra}} \in \{0.1, 0.2, 0.5, 1.0, 2.0\}$) shows the advantage peaks near $\gamma_{\text{extra}} = 0.5$–$1.0$ and collapses at very high rates, confirming that the paper's choice $\gamma_{\text{extra}} = 0.5$ is near-optimal. Taken together, local number variance functions as a finite-size, regime-dependent redistribution handle in the positive pocket ($J/U \geq 0.30$): it identifies intervention subsets with enhanced finite-horizon response, and this advantage is not reducible to geometric centrality, shell position, mean occupation, or disorder amplitude, while surviving the tested deterministic symmetry-breaking controls. Because the systems are small enough for exhaustive enumeration, we further rank the F_i-selected subset against all $\binom{L}{k}$ possible intervention subsets: in the positive-pocket regime it lies at the top of the subset-response distribution at every tested (L, τ) pair, and the conclusion survives signed, absolute, and redistribution-based target definitions. Finite-difference susceptibility analysis shows that F_i is not a local-depletion susceptibility but a redistribution susceptibility: Spearman$(F_i, \chi^{\text{redist}})$ is positive across all regimes and all horizons, and the top-F_i sites exactly match the highest redistribution-susceptibility sites at $L = 6$, $\tau \geq 2$.

I. INTRODUCTION

The one-dimensional Bose–Hubbard model is a standard setting for studying the interplay of quantum tunnelling and on-site interactions in interacting boson systems [? ?]. At zero temperature the model exhibits a quantum phase transition between a superfluid phase (large J/U) and a Mott-insulating phase (small J/U), controlled by the ratio of tunnelling amplitude J to on-site interaction U [? ?]. The clean-system phase diagram and its experimental realisation in optical lattices [?] have been extensively studied [?].

In the presence of dissipation—modelled by local dephasing in the Lindblad–Gorini–Kossakowski–Sudarshan (LGKS) framework [? ? ?]—the system is driven away from the ground state and develops spatially heterogeneous fluctuation structure. Open quantum systems of this type arise naturally in cold-atom experiments where coupling to photon-field fluctuations causes local dephasing of on-site number states [? ?]. More broadly, the engineering of dissipation as a resource for state preparation and quantum control has attracted significant theoretical and experimental attention [? ? ? ?].

A natural question in any driven-dissipative system is whether local structural variables can serve as targeted control handles: does selectively perturbing sites identified by a local observable produce measurably different future outcomes than applying the same perturbation budget at randomly chosen sites? In open quantum systems governed by a Lindblad master equation this question has received little direct quantitative attention.

This paper tests one specific instance. Under uniform dephasing, sites in the Bose–Hubbard chain develop different local occupation variances $F_i = \langle \hat{n}_i^2 \rangle - \langle \hat{n}_i \rangle^2$, reflecting the degree of local number fluctuation. We ask whether selecting the highest-F_i sites for additional dephasing produces a different pattern of future occupa-

* kunal29bhatia@gmail.com
† ORCID: 0009-0007-4447-6325

tion loss than applying the same total dephasing budget to randomly chosen sites. The question is about the interventional accessibility of future dynamical outcomes through structurally targeted dephasing, not about the equilibrium phase diagram. Here "handle" is used operationally: a site selector is called a redistribution handle if applying an explicit additional Lindblad dephasing term at the selected sites produces a statistically distinguishable future response from a matched-budget random-site intervention, with the pre-intervention state held fixed. This is an intervention contrast inside a fully specified dynamical model, not statistical causal identification from observational data.

In the clean, reflection-symmetric chain the high-F_i sites coincide with the geometrically central sites, so the primary sweep alone cannot attribute the effect to variance rather than geometry. We therefore perform four additional experiments designed to break this coincidence: a disorder selector sweep comparing F_i against eight alternative selectors across 50 disorder realizations, shell-matched permutation controls that hold shell geometry fixed while permuting sites within shells, a deterministic inhomogeneous (tilt) chain where $F_i \neq$ geo by construction, and a scan of the extra dephasing rate. Within the tested scope these controls resolve the clean-chain coincidence in favour of F_i: the variance advantage is not reducible to geometric centrality, shell position, mean occupation, disorder amplitude, or generator-action-based selection.

II. MODEL AND METHODS

A. Bose–Hubbard Hamiltonian

The Hamiltonian for a chain of L sites with open boundary conditions is

$$\hat{H} = -J \sum_{i=1}^{L-1} \left(\hat{a}_i^\dagger \hat{a}_{i+1} + \text{h.c.}\right) + \frac{U}{2} \sum_{i=1}^{L} \hat{n}_i(\hat{n}_i - 1), \quad (1)$$

where J is the tunnelling amplitude, U the on-site interaction, $\hat{n}_i = \hat{a}_i^\dagger \hat{a}_i$, and the local Fock space is truncated at $n_{\max} = 3$ occupations per site. The system is studied at half-filling ($N = \lfloor L/2 \rfloor$ particles) in the fixed-particle-number sector. Open boundary conditions are used throughout, so that the local variance field F_i can develop nontrivial spatial structure across the chain.

B. Hilbert space

The computation is restricted to the sector of exactly N particles distributed across L sites, each with at most $n_{\max} = 3$ occupations. The dimension of this sector is $\binom{L+N-1}{N}$ reduced by the $n_{\max}$ constraint; for $L = 6$, $N = 3$ the dimension is $D = 56$; for $L = 7$, $N = 3$ it is $D = 84$; for $L = 8$, $N = 4$ it is $D = 322$. For the two smaller sizes ($D \leq 84$), all operators are constructed as dense matrices and time evolution is computed via the direct Padé matrix-exponential path. For the largest tested size ($D = 322$), the Liouvillian superoperator (acting on the D^2-dimensional vectorised density matrix) is stored as a sparse CSR matrix; time evolution uses the Krylov-based `expm_multiply` path, and the per-trial modified Liouvillian is handled as a lightweight linear operator that avoids allocating a new sparse matrix per trial. The $n_{\max} = 3$ cutoff is binding only at $L = 8$, $N = 4$; a truncation check repeating those conditions at $n_{\max} = 4$ ($D = 322 \to 330$) left the regime ordering, sign pattern, and top-F_i subset-percentile tier unchanged, with local-loss gap shifts below 2×10^{-4} (see Discussion).

C. Lindblad dynamics

The density matrix $\hat{\rho}$ evolves under the Lindblad master equation

$$\frac{d\hat{\rho}}{dt} = -i[\hat{H}, \hat{\rho}] + \sum_{i=1}^{L} \gamma_i \Big(\hat{L}_i \hat{\rho} \hat{L}_i^\dagger - \tfrac{1}{2}\{\hat{L}_i^\dagger \hat{L}_i,\, \hat{\rho}\}\Big), \quad (2)$$

where $\hat{L}_i = \hat{n}_i$ is a dephasing jump operator and γ_i is the dephasing rate at site i. Baseline dephasing is uniform: $\gamma_i = \gamma = 0.1$ for all sites. Since $\hat{L}_i = \hat{n}_i$ commutes with the total number operator $\hat{N} = \sum_i \hat{n}_i$, the dissipator decoheres number-state superpositions without changing any local expectation value $\langle \hat{n}_i \rangle$; both local and total particle number are conserved at all times.

The Lindblad equation is converted to a superoperator acting on the vectorised density matrix in the row-major convention. Time evolution is computed by the action of the matrix exponential on the state vector, using `scipy.sparse.linalg.expm_multiply` [**?**], which avoids forming the full propagator. No Trotter decomposition, stochastic unravelling, or other approximation is used.

D. Burn-in state

The initial state is the ground state of $\hat{H}$, obtained by exact diagonalisation. This pure state is then evolved under uniform baseline dephasing for a burn-in period $t_{\text{burn}} = 5.0$ (in units of $\hbar/U$). The resulting mixed state $\hat{\rho}_{\text{burn}}$ serves as the starting point for all interventions within a given condition. The burn-in ensures that the state has developed spatially heterogeneous fluctuation structure under dissipation before any intervention is applied.

E. Handle definition

The handle variable is the local occupation variance

$$F_i = \langle \hat{n}_i^2 \rangle - \langle \hat{n}_i \rangle^2, \quad (3)$$

evaluated on the post-burn-in state $\hat{\rho}_{\text{burn}}$. This quantity identifies sites with larger number fluctuations. Because the Hamiltonian, the ground state, and the baseline Lindblad operators are all invariant under site reflection $i \leftrightarrow L-1-i$, the density matrix $\hat{\rho}(t)$ remains symmetric at all times, and consequently $F_i(t) = F_{L-1-i}(t)$ exactly. The F_i profile therefore peaks at the geometrically central sites and is smallest at the edges at every point in the evolution, including at t_{burn}. The top-k sites by F_i ($k = \lceil L/3 \rceil$) are selected for intervention. The selected set is fixed by the post-burn-in state and is therefore deterministic within each condition.

F. Intervention protocol

The **targeted intervention** applies additional dephasing $\gamma_{\text{extra}} = 0.5$ at the selected high-F_i sites, keeping all other sites at the baseline rate γ. The **random control** applies the same total extra budget ($k \times \gamma_{\text{extra}}$) to k randomly chosen sites. In both cases the total dissipative budget is identical. Because $\hat{L}_i = \hat{n}_i$ commutes with the total number operator, $\langle \hat{n}_i \rangle$ is unchanged at the moment of intervention by construction; the occupation profile is held fixed at the start of each arm.

The targeted intervention is deterministic: for a given condition, the post-burn-in state and the selected sites are the same, so the targeted evolution is computed once per (condition, τ) pair. Only the random arm varies across trials.

G. Target

The primary target is local future occupation loss at the intervention sites:

$$\mathcal{L}_{\text{local}}(\tau) = \sum_{i \in \mathcal{S}} \max\big(0,\ \langle \hat{n}_i \rangle_t - \langle \hat{n}_i \rangle_{t+\tau}\big), \tag{4}$$

where $\mathcal{S}$ is the set of selected sites and τ is the measurement horizon. The $\max(\cdot, 0)$ clip counts only sites that *lose* occupation; gains at other sites do not offset local losses. The clipped metric is the primary operational target because the intervention question asks whether selected sites subsequently lose occupation under targeted dephasing. Since dephasing conserves total particle number, local loss at selected sites necessarily reflects redistribution rather than particle removal. Signed, absolute, and redistribution-based metrics are reported in Sec. **??** as robustness and mechanistic checks.

H. Geometry–variance separation

Because the clean chain's reflection symmetry makes the high-F_i selector and a geometric central-k selector identical, the paper explicitly tests their disentanglement through three complementary controls: (i) random on-site disorder that breaks reflection symmetry across many realizations, allowing a nine-way selector comparison; (ii) shell-matched permutation controls that hold shell membership fixed and permute site assignment within each shell; and (iii) a deterministic inhomogeneous tilt that assigns site-dependent chemical potentials so that $F_i \neq$ geo by construction. Together these controls establish that the advantage of F_i targeting in the positive pocket is not an artefact of spatial centrality.

I. Design and statistical analysis

Four tunnelling-to-interaction ratios $J/U \in \{0.12, 0.20, 0.30, 0.40\}$ are tested across all chain lengths, spanning from deep Mott-like to intermediate superfluid-like regimes. To resolve the crossover boundary, a secondary boundary sweep tests $J/U \in \{0.16, 0.18, 0.20, 0.22, 0.24, 0.26, 0.28, 0.30\}$ at $L = 8$, using horizons $\tau \in \{2, 3\}$. Three horizons $\tau \in \{1, 2, 3\}$ and three chain lengths $L \in \{6, 7, 8\}$ are used in the primary sweep. For each condition, $N_{\text{trials}} = 100$ independent random-site draws are compared against the single targeted intervention. The test statistic is the mean difference in $\mathcal{L}_{\text{local}}$ between targeted and random interventions. A 95% bootstrap confidence interval (1000 resamples, trial-level resampling) is constructed for each condition. The decision rule is that a condition passes if the lower bound of the 95% CI exceeds zero. Because the targeted arm is deterministic for a fixed condition (the post-burn-in state and the F_i selector are both fixed), the bootstrap CI reflects trial-level variability in the matched random-control arm, not uncertainty over independent targeted realizations. The interval should therefore be read as uncertainty in the targeted-minus-random intervention contrast induced by random-site selection, not as a standard two-sample confidence interval. Bootstrap CIs are used descriptively throughout. Because multiple conditions are inspected, isolated CI crossings are not interpreted as standalone discoveries; the main claims rest on structured regime consistency across L and τ, exhaustive subset enumeration (Sec. **??**), and target-definition robustness. Random controls sample k-site subsets uniformly at random without replacement within each trial; trials are independent draws, and duplicate subsets across trials are not excluded. F_i ties in reflection-symmetric chains are resolved deterministically by the fixed argsort ordering; exhaustive enumeration confirms the positive-pocket claim is not dependent on random-baseline sampling.

J. Numerical diagnostics

Table **??** reports numerical consistency checks for representative conditions spanning the dense ($D = 56$) and sparse ($D = 322$) code paths, evaluated at $\tau = 3$ (the

TABLE I. Numerical diagnostics at $\tau = 3$ for representative conditions. $\varepsilon_{\rm tr} = |{\rm Tr}\,\rho_\tau - 1|$; $\varepsilon_{\rm H} = \|\rho_\tau - \rho_\tau^\dagger\|_\infty$; $\varepsilon_N = |\langle\hat{N}\rangle_\tau - \langle\hat{N}\rangle_0|$; $\lambda_{\min}$ is the smallest eigenvalue of ρ_τ. All errors are at double-precision machine level.

L	J/U	D	$\varepsilon_{\rm tr}$	$\varepsilon_{\rm H}$	ε_N	$\lambda_{\min}$
6	0.12	56	4×10^{-16}	3×10^{-17}	4×10^{-16}	4×10^{-7}
6	0.40	56	4×10^{-16}	2×10^{-17}	4×10^{-16}	3×10^{-4}
8	0.12	322	7×10^{-16}	6×10^{-18}	1×10^{-15}	2×10^{-10}
8	0.40	322	2×10^{-16}	4×10^{-18}	4×10^{-16}	8×10^{-6}

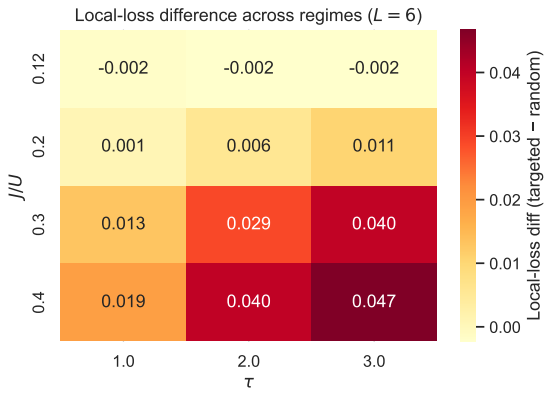

FIG. 1. Local-loss difference across J/U and τ at $L = 6$. The handle effect is positive in the intermediate and stronger coupling regime and absent or reversed at the lowest tested coupling.

longest tested horizon). These are not additional physical assumptions; they document the floating-point residuals of the exact `expm_multiply` evolution.

III. RESULTS

A. Primary sweep: clean chain, $L \in \{6, 7, 8\}$

Table **??** reports the local-loss difference for all conditions at $L = 6$. At $J/U = 0.12$ the response is near-zero to weakly harmful: the 95% CI is below zero at the paper's burn-in time, but the gaps are small in magnitude ($|\text{mean}| \lesssim 0.01$) and, as shown in Sec. **??**, partly burn-in-sensitive at $L = 6$. At $J/U = 0.20$ there is a crossover that is sensitive to system size and horizon. At $J/U \geq 0.30$ the targeted intervention is reliably *beneficial*, with CIs strictly above zero at every tested (L, τ) combination. The largest effects occur at $J/U = 0.40$, $\tau = 3$, growing with system size: $+0.047$ ($L = 6$), $+0.068$ ($L = 7$), $+0.072$ ($L = 8$).

The effect increases with J/U from the negative boundary at 0.12 through a transitional regime at 0.20 and into the robust positive pocket at 0.30–0.40 (Fig. **??**). The magnitude of the positive difference also increases with the time horizon τ, consistent with a cumulative dynamical effect rather than a transient perturbation.

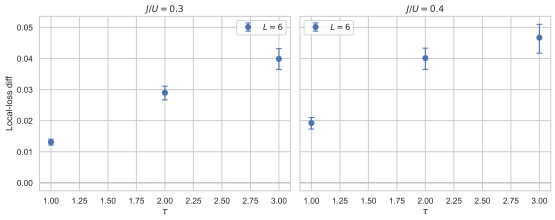

FIG. 2. Local-loss difference at $J/U = 0.30$ and 0.40 for $L = 6$ and $L = 7$ (two representative sizes). Error bars are 95% bootstrap CIs. The same positive result holds at $L = 8$; the positive pocket is robust across all tested $L \in \{6, 7, 8\}$, with the effect growing with L at the strongest coupling.

B. Robustness across chain length

Table **??** shows pooled results at $J/U \in \{0.30, 0.40\}$ for $L \in \{6, 7\}$. The positive pocket at $J/U \in \{0.30, 0.40\}$ passes at all three tested chain lengths $L \in \{6, 7, 8\}$ with 95% confidence intervals excluding zero at every horizon $\tau \in \{1, 2, 3\}$ (Fig. **??**), confirming that the positive pocket is not a small-size artefact.

The intermediate coupling $J/U = 0.20$ is strongly size- and horizon-sensitive, placing it in the crossover band rather than the robust positive pocket (see Sec. **??**).

C. Regime map

The results define three regimes within the tested parameter range:

Low-coupling regime ($J/U = 0.12$): The targeted intervention produces near-zero to weakly negative responses relative to random targeting. In this deep Mott-like regime the occupation-variance handle does not identify intervention-sensitive sites; the 95% CIs are below zero at the paper's burn-in time, but the gaps are small in magnitude and, as shown in Sec. **??**, partly burn-in-sensitive at $L = 6$.

Crossover band ($J/U \approx 0.18$–0.24, depending on L and τ): The handle begins to function but its effectiveness is size- and horizon-dependent. The first-pass onset shifts upward with L (for fixed τ) and downward with τ (for fixed L), both trends monotone across the ladder (see Sec. **??**).

Positive pocket ($J/U \gtrsim 0.30$): The handle is consistently positive across all three tested chain lengths and all horizons $\tau \in \{1, 2, 3\}$. This is the regime where the occupation-variance selector reliably identifies sites whose perturbation produces excess local loss, with CIs strictly above zero at every tested (L, τ) combination.

The existence of a negative boundary regime and a structured crossover band is itself informative: the handle is not a generic property of the dephasing operator

TABLE II. Local-loss difference (targeted − random) at $L = 6$, $N_{\text{trials}} = 100$ per condition. Format: mean [95% CI]. $J/U = 0.12$ is near-zero to weakly negative (low-coupling regime); $J/U = 0.20$ is transitional (crossover band); $J/U \in \{0.30, 0.40\}$ are robustly positive.

J/U	$\tau = 1$	$\tau = 2$	$\tau = 3$
0.12	−0.002 [−0.002, −0.001]	−0.002 [−0.003, −0.002]	−0.002 [−0.003, −0.001]
0.20	0.001 [0.000, 0.001]	0.006 [0.005, 0.007]	0.011 [0.009, 0.012]
0.30	0.013 [0.012, 0.014]	0.029 [0.027, 0.031]	0.040 [0.036, 0.043]
0.40	0.019 [0.017, 0.021]	0.040 [0.036, 0.043]	0.047 [0.042, 0.051]

TABLE III. Pooled local-loss difference for $L \in \{6, 7\}$ at $J/U = 0.30$ and 0.40 (two representative sizes from the full ladder). Format: mean [95% CI], pooling trial-level differences across $\tau \in \{1, 2, 3\}$. All three tested sizes $L \in \{6, 7, 8\}$ pass at both couplings (see text).

Condition	$L = 6$	$L = 7$
$J/U = 0.30$	0.027 [0.025, 0.029]	0.030 [0.027, 0.032]
$J/U = 0.40$	0.035 [0.032, 0.038]	0.047 [0.044, 0.051]

TABLE IV. First-PASS onset coupling: smallest J/U at which the 95% CI strictly excludes zero, by chain length L and horizon τ. Entries ≤ 0.20 are bounded by the initial sweep (exact onset lies in the interval (0.12, 0.20]); entry > 0.20 indicates a failure at $J/U = 0.20$ with onset in (0.20, 0.30).

L	$\tau = 2$	$\tau = 3$
6	≤ 0.20	≤ 0.20
7	> 0.20	≤ 0.20
8	0.20	0.18

but is sharply regime-dependent.

D. Crossover-band onset

Table **??** reports the smallest J/U at which the 95% CI first strictly excludes zero (first PASS) for $L \in \{6, 7, 8\}$ and horizons $\tau \in \{2, 3\}$.

The onset coupling is non-decreasing with L for fixed τ, and non-increasing with τ for fixed L, across the full tested range. Any coupling $J/U \gtrsim 0.30$ lies above the onset for all tested (L, τ) combinations, consistently defining the positive pocket boundary.

E. Spatial response pattern

Fig. **??** shows the site-resolved occupation change at $L = 6$, $\tau = 2$ for both regimes side by side: the low-coupling regime (panel a, $J/U = 0.12$) and the positive pocket (panel b, $J/U = 0.30$). The bars show the mean site-resolved occupation change under the deterministic targeted intervention and averaged across the 100 matched-budget random interventions.

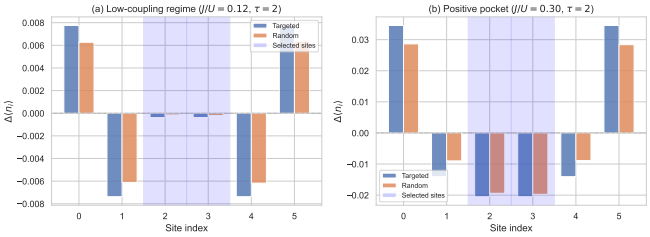

FIG. 3. Mean site-resolved occupation change ($\Delta\langle n_i \rangle$, $L = 6$, $\tau = 2$) under targeted versus matched-budget random intervention, shown for both regimes. **(a)** Low-coupling regime ($J/U = 0.12$): high-F_i selected sites (shaded, sites 2–3) show near-zero response; the large occupation redistribution occurs at unselected sites 1 and 4, where targeting provides no advantage. **(b)** Positive pocket ($J/U = 0.30$): targeted intervention drives deeper occupation redistribution than random intervention across the non-selected sites, consistent with the redistribution-handle interpretation. Negative values correspond to occupation loss; shaded bands mark selected sites.

The positive-pocket panel (b) is consistent with a nonlocal, redistributive response: under targeted intervention, the selected central sites (shaded) show less negative occupation change than under random intervention, while non-selected sites show increased occupation loss, indicating that the intervention redistributes occupation across the chain rather than simply amplifying local loss at the targeted sites. The low-coupling panel (a) shows the regime-dependent contrast directly: at $J/U = 0.12$, the high-F_i selected sites (2,3) carry near-zero occupation change, while sites 1 and 4—outside the selected set—carry the dominant redistribution signal. Targeting the variance handle therefore perturbs the same structural degree of freedom in both regimes, but whether that perturbation leads to net local loss at the selected sites or redistributes occupation in a non-beneficial direction depends on the coupling strength.

F. Disorder selector sweep

To test whether the advantage of F_i targeting is specific to variance information or reducible to geometry, we

TABLE V. Selector definitions used in the disorder sweep. Each selector picks the top-k (or bottom-k) sites by the listed criterion evaluated on the post-burn-in state ρ_t.

Selector	Definition
fi	Top-k by $F_i = \langle \hat{n}_i^2 \rangle - \langle \hat{n}_i \rangle^2$
geo	k sites nearest the chain centre
maxn	Top-k by $\langle \hat{n}_i \rangle$
minn	Bottom-k by $\langle \hat{n}_i \rangle$
bdy	k outermost (boundary) sites
anti	Bottom-k by F_i
gen	Top-k by $\lvert\langle [\hat{L}_i, \hat{H}] \rangle\rvert$, generator action
dis_amp	Top-k by $\lvert\mu_i\rvert$, disorder amplitude
dis_anti	Bottom-k by $\lvert\mu_i\rvert$

broke the clean-chain reflection symmetry by adding random on-site disorder $\mu_i \sim \mathrm{Uniform}(0, \mu_{\max})$ to the chemical potential. Nine selectors were compared at $L = 6$, 50 disorder realizations, and $J/U = 0.40$, $\tau = 3$. Their definitions are given in Table **??**.

At strong disorder ($\mu_{\max} \geq 1.0$) fi is the top selector (Fig. **??**, left): fi= +0.049, maxn= +0.034, geo= +0.032, gen= +0.028, dis_amp= +0.023, bdy= +0.010, dis_anti= +0.005, anti= +0.002, minn≈ 0. The fi–geo gap grows monotonically with both $\mu_{\max}$ and J/U, from +0.0002 at $\mu_{\max} = 0.5$ to +0.017 at $\mu_{\max} = 2.0$. The fact that dis_amp ranks well below fi confirms that F_i carries information beyond raw disorder amplitude: variance tracks local susceptibility to dephasing, not merely where disorder is strongest. At weak disorder ($\mu_{\max} \lesssim 0.5$), symmetry is approximately intact and fi≈geo, consistent with the clean-chain result.

G. Shell-matched permutation controls

In the disordered chain, sites can be grouped into shells by their distance from the nearest boundary. Shell-matched permutation controls hold shell membership fixed and randomly permute the assignment of sites within each shell, then apply targeted dephasing to the permuted top-k sites. This isolates the within-shell variance information from the shell-level geometry.

At $L = 6$, $\mu_{\max} \leq 0.3$, $J/U \geq 0.30$, $\tau = 3$, fi beats all shell-matched permutations by +0.002 to +0.006 (Fig. **??**, right). This establishes that F_i carries within-shell information that is not captured by shell geometry alone: sites sharing the same shell distance are not interchangeable when selected by variance.

H. Deterministic inhomogeneous chain

To test the geometry–variance separation without stochasticity, we studied a deterministic tilt chain in

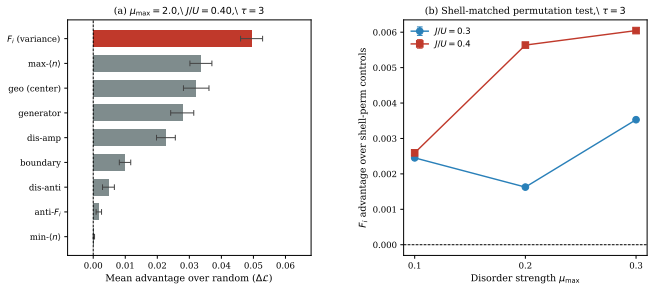

FIG. 4. Selector discrimination under symmetry breaking. **Left:** under strong disorder, the variance-based selector F_i ranks above the tested alternatives, including geometric centrality and disorder amplitude. **Right:** shell-matched permutation controls show that F_i retains an advantage beyond shell position alone. Together these results show that strong disorder separates variance information from geometry and that the surviving advantage is not reducible to shell membership.

which a linearly varying chemical potential $\mu_i = \mu_0(i - (L-1)/2)$ is applied, explicitly breaking reflection symmetry by design. The resulting F_i profile does not peak at the geometric center, so fi and geo select different sites (overlap zero).

At $J/U = 0.40$, $\tau = 3$: $F_i >$ geo with a gap of +0.032 to +0.054 (Fig. **??**, left). This non-random, symmetry-breaking control confirms that variance tracks local number fluctuation structure rather than spatial position, and that the handle is not simply a proxy for geometric centrality.

I. Extra dephasing rate scan

To assess sensitivity to the choice of γ_{extra}, we scanned $\gamma_{\mathrm{extra}} \in \{0.1, 0.2, 0.5, 1.0, 2.0\}$ at $L = 6$, $J/U = 0.40$, $\tau = 3$. The mean advantages are +0.0063, +0.0106, +0.0163, +0.0149, +0.0035 respectively. The advantage peaks near $\gamma_{\mathrm{extra}} = 0.5$–$1.0$ and collapses at very high rates, where the extra dephasing saturates the dynamics (Fig. **??**, right). The paper's primary choice $\gamma_{\mathrm{extra}} = 0.5$ lies at the near-optimal point rather than at an arbitrary or cherry-picked value.

J. Exhaustive subset ranking and target robustness

The comparison against 100 sampled random controls establishes that the F_i-selected subset outperforms a matched-budget random baseline, but does not bound its rank among *all* possible subsets. Because the Hilbert spaces studied here are small enough for exhaustive enumeration, we evaluated every $\binom{L}{k}$ possible k-site subset ($k = \lceil L/3 \rceil$; $\binom{6}{2} = 15$, $\binom{7}{3} = 35$, $\binom{8}{3} = 56$ subsets) under the paper's exact Lindblad protocol, and ranked the F_i-selected subset by its clipped local-loss percentile.

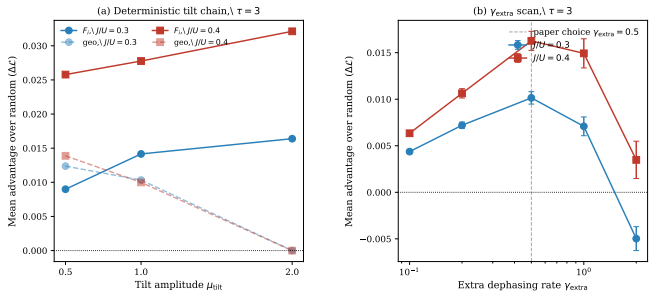

FIG. 5. Symmetry breaking and intervention-strength robustness. **Left:** in a deterministic tilted chain where the variance and geometric selectors differ by construction, the variance-based selector outperforms geometric centrality. **Right:** the intervention advantage is robust across a broad mid-range of extra dephasing strength and is near-optimal around $\gamma_{\text{extra}} \approx 0.5$–$1.0$, showing that the main paper choice is not a tuning artefact.

TABLE VI. Percentile rank of the F_i-selected subset among all $\binom{L}{k}$ intervention subsets, clipped local-loss metric. 100 = globally optimal (tied or unique).

L	J/U	$\tau = 1$	$\tau = 2$	$\tau = 3$
6	0.30	100	100	100
6	0.40	100	100	100
7	0.30	100	100	100
7	0.40	100	100	100
8	0.30	100	100	100
8	0.40	98	98	98
6	0.20	80	87	87
7	0.20	29	46	66
8	0.20	34	68	71
6	0.12	40	40	40
7	0.12	29	29	29
8	0.12	21	21	21

Table **??** reports the results. In the positive-pocket regime ($J/U \geq 0.30$), the F_i-selected subset is globally optimal or tied for optimal in all $L \in \{6, 7\}$ conditions and at $L = 8$, $J/U = 0.30$, and remains in the 98th percentile at $L = 8$, $J/U = 0.40$ (one of 56 subsets ties or exceeds it). The crossover band ($J/U = 0.20$) yields 29–87th percentile depending on L and τ, consistent with the regime's size- and horizon-sensitive character. The low-coupling regime ($J/U = 0.12$) gives 21–40th percentile—the F_i-selected subset actively ranks below the median, consistent with the low-coupling-regime result.

We also verified that the positive-pocket advantage is not produced by the positive-part clipping in the target metric. For every (L, τ) condition with $J/U \geq 0.30$: the signed local-change gap (no clipping floor) is positive and exceeds the clipped gap; the absolute local-change gap is positive; and the global redistribution gap $\sum_j |\langle \hat{n}_j \rangle_{\text{targeted}} - \langle \hat{n}_j \rangle_{\text{baseline}}|$ is positive. The clipping is conservative, not inflating: removing it yields a *larger* advantage in every tested positive-pocket condition.

TABLE VII. Burn-in sensitivity at $\tau = 3$: clipped local-loss gap and exhaustive-subset percentile at three burn-in multipliers relative to $t_{\text{burn}} = 5$. Positive gap means the F_i-selected subset exceeds the all-subset mean.

L	J/U	0.5× gap	1.0× gap	2.0× gap	Pct. range	Sign stable
6	0.12	−0.002	−0.002	+0.003	40–80%	no
6	0.40	+0.064	+0.047	+0.012	87–100%	yes
8	0.12	−0.005	−0.006	< 0.001	21–54%	yes
8	0.40	+0.065	+0.073	+0.009	80–100%	yes

K. Burn-in sensitivity

Table **??** reports the clipped local-loss gap and exhaustive subset percentile at $\tau = 3$ for three burn-in multipliers (0.5×, 1.0×, 2.0× the paper's baseline $t_{\text{burn}} = 5$). The positive-pocket conclusion ($J/U = 0.40$) is stable at both $L = 6$ and $L = 8$: the gap remains positive and the F_i-selected subset stays in the top quartile or above across all three burn-in times. The low-coupling response ($J/U = 0.12$) is much smaller in magnitude ($|\text{gap}| \lesssim 0.006$) and partly burn-in-sensitive at $L = 6$. The positive-pocket conclusion is stable under 0.5×–2× burn-in variation at both $L = 6$ and $L = 8$; the low-coupling response is much smaller in magnitude and can become burn-in-sensitive at $L = 6$.

L. Susceptibility structure: redistribution, not local depletion

To probe whether F_i is grounded in the actual per-site dephasing response, we compute finite-difference susceptibilities at each site i under a small extra dephasing dose $\delta\gamma = 0.1$:

$$\chi_i^{\text{clip}}(\tau) = \frac{\max(0,\ \langle \hat{n}_i \rangle_0^\tau - \langle \hat{n}_i \rangle_{\delta\gamma_i}^\tau)}{\delta\gamma}, \tag{5}$$

$$\chi_i^{\text{redist}}(\tau) = \frac{\sum_j |\langle \hat{n}_j \rangle_{\delta\gamma_i}^\tau - \langle \hat{n}_j \rangle_0^\tau|}{\delta\gamma}, \tag{6}$$

where $\langle \hat{n}_j \rangle_{\delta\gamma_i}^\tau$ is the occupation at site j after evolving to horizon τ with extra dephasing $\delta\gamma$ applied only at site i, and $\langle \hat{n}_j \rangle_0^\tau$ is the corresponding baseline (no extra dephasing).

The Spearman correlation between F_i and these susceptibilities, computed over all L sites for each $(L, J/U, \tau)$ condition, reveals a clear structural distinction. The sign pattern is stable across finite-difference doses $\delta\gamma \in \{0.05, 0.1, 0.2\}$. F_i is *not* a local-depletion susceptibility: Spearman$(F_i, \chi^{\text{clip}})$ is negative at $\tau = 1, 2$ in nearly all conditions and only weakly positive at $\tau = 3$ in the

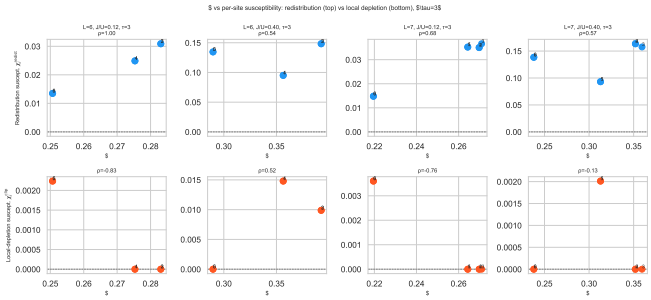

FIG. 6. Per-site susceptibility vs F_i at $\tau = 3$ for representative conditions. **Top row:** redistribution susceptibility χ_i^{redist} vs F_i — Spearman correlation is positive in all four panels. **Bottom row:** local-depletion susceptibility χ_i^{clip} vs F_i — correlation is negative or near zero in three of four panels, and weakly positive only at $J/U = 0.40$. The contrast establishes that F_i is a redistribution handle, not a local-depletion handle.

positive pocket. By contrast, F_i is consistently a *redistribution* susceptibility: Spearman$(F_i, \chi^{\text{redist}})$ is positive across all regimes and all τ, reaching values of 0.5–1.0 at $\tau \geq 2$ (Fig. **??**). At $L = 6$, $\tau \geq 2$, the top-k overlap between the highest-F_i sites and the highest-χ^{redist} sites is 1.0 in all tested conditions, meaning F_i exactly identifies the sites that most strongly perturb the global occupation distribution per unit extra dephasing.

This susceptibility structure provides a mechanistic interpretation of the regime-dependence: F_i consistently identifies sites with high redistribution susceptibility, but whether that redistribution leads to net local occupation loss at the selected sites (positive pocket) or net local gain (negative regime) depends on the global response pattern at that coupling strength. The variance handle works not by directly depleting targeted sites but by controlling how occupation is redistributed across the chain on the time scale τ.

IV. DISCUSSION

Summary. In the tested one-dimensional open Bose–Hubbard chain, local number variance $F_i = \langle \hat{n}_i^2 \rangle - \langle \hat{n}_i \rangle^2$ functions as a finite-size, regime-dependent redistribution handle: in the positive pocket $J/U \geq 0.30$, targeting the highest-F_i sites with additional dephasing produces more local future occupation loss than matched-budget random targeting, robustly across all tested $L \in \{6, 7, 8\}$ and $\tau \in \{1, 2, 3\}$, with the effect increasing with system size. Exhaustive enumeration of all $\binom{L}{k}$ intervention subsets further shows that the F_i-selected subset lies at the top of the subset-response distribution in the positive pocket, and this advantage survives signed, absolute, and redistribution-based target definitions. The handle is regime-dependent—near-zero to weakly harmful at $J/U = 0.12$, transitional near $J/U \approx 0.20$—and its advantage is not reducible to geometric centrality, shell position, mean occupation, or disorder amplitude.

Mechanism. The susceptibility analysis in Sec. **??** shows that F_i is a redistribution susceptibility rather than a local-depletion susceptibility: high-F_i sites are those where a small extra dephasing dose most strongly perturbs the global occupation pattern. Whether this redistribution results in net local loss at the targeted sites (positive pocket) or net local gain (low-coupling regime) depends on the coupling-strength-dependent global response structure, not on any local property of the sites themselves. A full microscopic account of this redistribution—connecting F_i to the Lindblad generator spectrum or the Green's function of the superoperator—is left for future work.

What is not claimed. The exhaustive enumeration establishes near-optimality within the finite-size systems studied; we do not claim this extends to the thermodynamic limit, to all coupling regimes, or to other quantum systems. No claim is made that $\gamma_{\text{extra}} = 0.5$ is a globally optimal dephasing rate, or that the occupation-variance handle corresponds to any equilibrium observable. The observed regime boundary near $J/U \approx 0.20$ should not be identified with the thermodynamic Mott–superfluid transition. In these small open chains under dissipation it is best interpreted as a finite-size dynamical crossover in the response of the Lindblad evolution, not as a sharp equilibrium phase boundary.

Scope and limitations. The Hilbert-space dimension of the fixed-N sector grows combinatorially with L and $n_{\max}$, limiting exact Lindblad evolution to small chains. The tested range ($L \leq 8$, $n_{\max} = 3$) establishes the effect in the positive pocket but does not access the thermodynamic limit. Whether the positive pocket persists, broadens, or shifts at larger system sizes is an open question requiring approximate methods such as matrix-product-operator Lindblad evolution [**?**]. As a truncation check, we repeated the $L = 8$, $N = 4$ conditions with $n_{\max} = 4$ (the only tested case where the $n_{\max} = 3$ cutoff is active, $D = 322 \to 330$); the regime ordering, sign pattern, and top-F_i subset-percentile tier are unchanged across all tested conditions, with local-loss gap shifts below 2×10^{-4}.

Experimental relevance. The 1D open Bose–Hubbard model under dephasing is experimentally motivated. Cold-atom optical-lattice platforms already realise 1D Bose–Hubbard-like systems with tunable J/U [**?** **?**], and recent advances in quantum-gas microscopy provide site-resolved preparation and readout relevant to the targeted-intervention protocol studied here [**?**]. The Lindblad dephasing modelled by $\gamma \hat{n}_i$ is a standard approximation for photon-scattering-induced decoherence in optical lattices [**?** **?**]. The present results therefore constitute a theory-side proposal with experimental plausibility: the positive-pocket and crossover-band structures identified here are, in principle, accessible to existing site-resolved cold-atom platforms, though the exact targeted-dephasing protocol has not yet been implemented experimentally.

ACKNOWLEDGMENTS

The author thanks the open-source scientific Python community for the libraries underlying this work: NumPy, SciPy [**?**], pandas, matplotlib, and tqdm.

DATA AND CODE AVAILABILITY

All simulation code and raw data are available at https://github.com/kunalb541/BH.